



Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

SSA Linear Algebra Final

Image Convolution in MATLAB (and a bit of Python)

Evan Xiang¹ Vivan Poddar¹ Rohan Khera¹ Beau Brush¹

¹Mathematics Department
Shady Side Academy Senior School

December 16th, 2024



Table of Contents

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

- 1 Overview
- 2 Gaussian
- 3 Cauchy
- 4 Box Blur
- 5 Kernel Implementation
- 6 Sliding Window
- 7 Interface
- 8 Proof
- 9 Future Work
- 10 Applications
- 11 Thanks!



Convolution Overview

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

Definition

A convolution is done by multiplying a pixel's and its neighboring pixels color value by a matrix.

Generalized Matrix Convolution Formula

$$(F * G)(x, y) = \sum_m \sum_n F(x - m, y - n) G(m, n)$$

Where:

- x and y represent the current position within the output matrix
- m and n are generalized variables representing the shift of the matrix with regard to $G(m, n) \rightarrow$ the kernel matrix



Convolution Cont.

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!




Figure: A general description of the convolution operation. (Photo downloaded from Apple Developer Documentation, CC BY-SA 4.0)

Summary of the Convolution Operation

- Place the kernel $G(m, n)$ over the input matrix $F(x, y)$ such that the kernel's center aligns with the current position (x, y) of the output matrix.
- For every element of the kernel $G(m, n)$, multiply it with the corresponding element of the input matrix $F(x - m, y - n)$.
- Sum all the resulting products to compute the value for the current position (x, y) in the output matrix.

$$(F * G)(x, y) = \sum_m \sum_n F(x - m, y - n) G(m, n)$$

- Slide the kernel over the input matrix to the next position and repeat steps 2 and 3. This involves systematically shifting the kernel across all valid positions in the input matrix.



Weierstrass Transform (Gaussian Filter)

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!




Figure: Weierstrass Transform for 5 parameters of t . The green line represents the standard Weierstrass. (Photo downloaded from Glosser.ca - Own work, CC BY-SA 4.0)

Definition

The Weierstrass transform applies a 2D Gaussian Kernel to an image to reduce noise and create a blurred version of the original function by taking a weighted average of the function's values with the degree of smoothing controlled by the Gaussian's variance.



Gaussian Cont.

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

Formula

$$O(i, j) = \sum_{x=-\lfloor N/2 \rfloor}^{\lfloor N/2 \rfloor} \sum_{y=-\lfloor N/2 \rfloor}^{\lfloor N/2 \rfloor} \left(\frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \right) \cdot I(i-x, j-y)$$

Where:

- N represents the size of the kernel
- $\left(\frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \right)$ represents the Gaussian kernel
- $O(i, j)$ represents the output



Cauchy Kernel

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

Cauchy Matrix Convolution Formula

$$O(i, j) = \sum_{x=-\lfloor N/2 \rfloor}^{\lfloor N/2 \rfloor} \sum_{y=-\lfloor N/2 \rfloor}^{\lfloor N/2 \rfloor} \frac{1}{\pi \gamma \left(1 + \frac{x^2 + y^2}{\gamma^2}\right)} \cdot I(i - x, j - y)$$

Note that this looks relatively similar to the Weierstrass Transform with the Gaussian! But how does it differ...

Gaussian vs. Cauchy Kernel

The primary difference arises in the difference in the kernel decay rate. The Gaussian, represented by $K_{\text{Gaussian}}(r) = \exp\left(-\frac{r^2}{2\sigma^2}\right)$, decreases exponentially. The Cauchy, represented by $K_{\text{Cauchy}}(r) = \frac{\gamma}{\pi(\gamma^2 + r^2)}$, decreases algebraically, allowing for longer distance influences.



Box Blur Method

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

Formula

$$O(i,j) = \frac{1}{N^2} \sum_{x=-\lfloor N/2 \rfloor}^{\lfloor N/2 \rfloor} \sum_{y=-\lfloor N/2 \rfloor}^{\lfloor N/2 \rfloor} I(i-x, j-y)$$

See how this is much simpler than all the previous algorithms... There's a reason for this... I mean it's literally just a weighted average of the values covered by the kernel applied to all values in the image matrix.



The Gaussian in Disguise

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

Proof

The box blur operation is a convolution of the image I with $K(x, y)$, producing a blurred output $O(x, y)$. Repeated application of the box blur corresponds to convolving $K(x, y)$ with itself n times. After n convolutions, the resulting kernel $K_n(x, y)$ is given by:

$$K_n(x, y) = K(x, y) * K(x, y) * \cdots * K(x, y) \quad (n \text{ times}).$$

By the Central Limit Theorem, the sum of n independent random variables converges to a Gaussian distribution as $n \rightarrow \infty$ with finite mean and variance. The kernel $K(x, y)$ acts as a probability distribution. Since the uniform box kernel $K(x, y)$ satisfies the conditions of the CLT, the repeated convolution $K_n(x, y)$ converges to a Gaussian kernel $G(x, y)$:

$$G(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right),$$

where the standard deviation σ grows proportionally to \sqrt{n} .

Thus, repeated application of the box blur approximates the Gaussian Kernel as $n \rightarrow \infty$.



Kernel Implementation

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

```
99 % kernel creation functions
100 function gaussianKernel = createGaussianKernel(kernelSize, sigma)
101     [x, y] = meshgrid(-floor(kernelSize/2):floor(kernelSize/2), -floor(kernelSize/2):floor(kernelSize/2));
102     gaussianKernel = exp(-(x.^2 + y.^2) / (2 * sigma^2));
103     gaussianKernel = gaussianKernel / sum(gaussianKernel(:));
104 end
105
106 function cauchyKernel = createCauchyKernel(kernelSize, gamma)
107     [x, y] = meshgrid(-floor(kernelSize/2):floor(kernelSize/2), -floor(kernelSize/2):floor(kernelSize/2));
108     cauchyKernel = 1 ./ (1 + (x.^2 + y.^2) / gamma^2);
109     cauchyKernel = cauchyKernel / sum(cauchyKernel(:));
110 end
111
112 function boxKernel = createBoxKernel(kernelSize)
113     boxKernel = ones(kernelSize, kernelSize) / (kernelSize^2);
114 end
115
```

Figure: Implementation of all 3 kernels in MATLAB (.mat) code. Each function defines a separate type of kernel which can be called to slide across the pixels of an image.



Sliding Window

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel

Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

```
197 for channel = 1:3
198     region = double(image(y:y+height-1, x:x+width-1, channel));
199     cauchyBlurredRegion = zeros(size(region));
200     boxBlurredRegion = zeros(size(region)); % Replaced laplacian with box
201     gaussianBlurredRegion = zeros(size(region));
202     [kHeight, kWidth] = size(cauchyKernel);
203     padHeight = floor(kHeight / 2);
204     padWidth = floor(kWidth / 2);
205     paddedRegion = padarray(region, [padHeight, padWidth], 0, 'both');
206     for i = 1:size(region, 1)
207         for j = 1:size(region, 2)
208             localWindow = paddedRegion(i:i+kHeight-1, j:j+kWidth-1);
209             cauchyBlurredRegion(i, j) = sum(sum(localWindow .* cauchyKernel));
210             boxBlurredRegion(i, j) = sum(sum(localWindow .* boxKernel)); % Replaced laplacian with box
211             gaussianBlurredRegion(i, j) = sum(sum(localWindow .* gaussianKernel));
212         end
213     end
214     cauchyBlurredImage(y:y+height-1, x:x+width-1, channel) = uint8(cauchyBlurredRegion);
215     boxBlurredImage(y:y+height-1, x:x+width-1, channel) = uint8(boxBlurredRegion); % Replaced laplacian with box
216     gaussianBlurredImage(y:y+height-1, x:x+width-1, channel) = uint8(gaussianBlurredRegion);
217 end
```

Figure: Implementation of the full convolution operation, using the 3 separate kernels defined in the previous slide.



Final Developed Interface

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

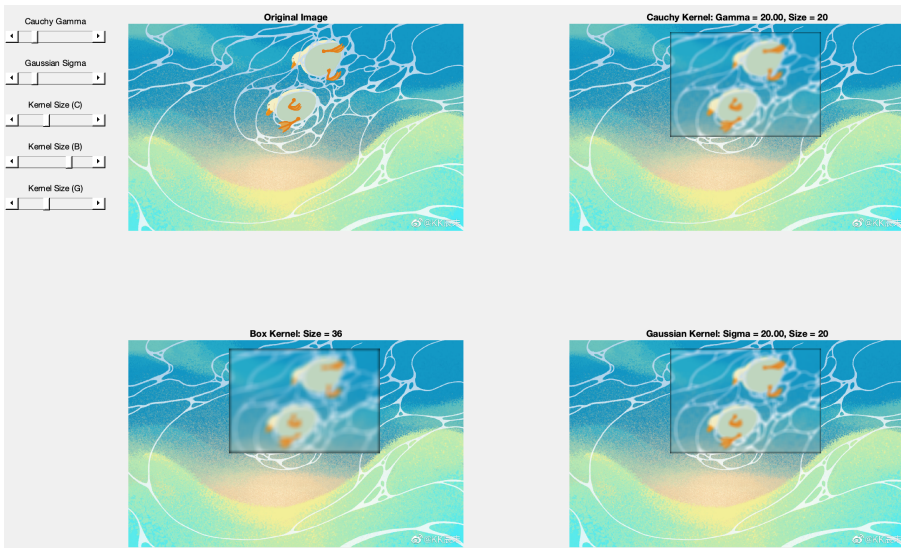
Interface

Proof

Future Work

Applications

Thanks!





Quod erat demonstrandum!

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel

Implementation

Sliding Window



Interface

Proof

Future Work

Applications

Thanks!





Code for Proof

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel

Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

```
✓ 3s ▶ import cv2
import numpy as np
import matplotlib.pyplot as plt
def gaus2d(x=0, y=0, mx=0, my=0, sx=1, sy=1):
    return 1. / (2. * np.pi * sx * sy) * np.exp(-((x - mx)**2. / (2. * sx**2.) + (y - my)**2. / (2. * sy**2.)))
image = np.ones((49, 49)) * 255
image[24][24] = 0
gblur = cv2.GaussianBlur(image, (31, 31), 3.5, 3.5)
box_blurs = [image]
for i in range(6):
    box_blurs.append(cv2.boxFilter(box_blurs[-1], -1, (5, 5)))
plt.figure(figsize=(18, 3))
titles = ["Input"] + [f"Box Blur - {i} Pass{'es' if i > 1 else ''}" for i in range(1, 7)] + ["Gaussian Blur"]
for i, (title, img) in enumerate(zip(titles, box_blurs + [gblur])):
    plt.subplot(1, len(titles), i + 1)
    plt.title(title)
    plt.imshow(img, cmap='Blues')
    plt.gca().get_xaxis().set_visible(False)
    plt.gca().get_yaxis().set_visible(False)
plt.tight_layout()
plt.savefig("box.jpg", facecolor='white', dpi=300)
plt.show()
```

Figure: Python code that displays the repeated box blurs after each iteration, in addition to the Gaussian blur.



Future Work

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

Future work would most likely involve the implementation of an algorithm to sharpen images using Laplacian of Gaussian.

Laplacian of Gaussian (LoG)

We apply the Laplacian operator directly to the Gaussian kernel, resulting in the following equation:

$$LoG(x, y) = \Delta (G(x, y) * I(x, y))$$

Then, we expand this with full equations to:

$$\Delta G(x, y) = \left(\frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} \right) \cdot \frac{1}{2\pi\sigma^2} \exp \left(-\frac{x^2 + y^2}{2\sigma^2} \right)$$

This equation lends itself to edge detection, as the Laplacian blurring function may selectively blur non-edges, while the Gaussian smooths out noise, resulting in edge-detection with LoG.



Convolutional Neural Networks (ConvNet)

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

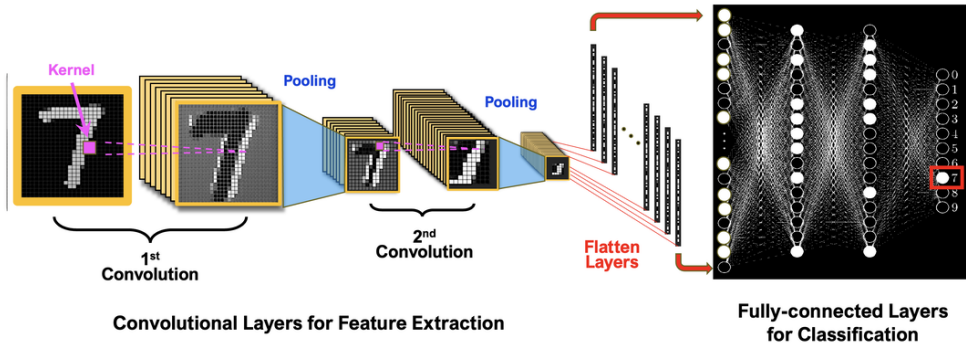
Interface

Proof

Future Work

Applications

Thanks!





Thanks for Listening

Linear Algebra

FINAL EXAM

Overview

Gaussian

Cauchy

Box Blur

Kernel
Implementation

Sliding Window

Interface

Proof

Future Work

Applications

Thanks!

- All information for implementation of our algorithm may be found here:
<https://github.com/evankxiang/ssalinalgfinal>
- We have included CSV files of the image convolution algorithm and all raw code for the algorithm (.mat files). Furthermore, we include our old test files, primarily consisting of non-selective image blurring (no ability to select a bounding box). Finally, we include a preliminary sharpening algorithm using LoG that is a WIP.
- EX, VP completed the programming segment. EX, RK, BB completed the slideshow and testing of the algorithms. Note, we are NOT listed in order of contribution.